

ECE 4750 PSET 2

Tim Yao (ty252)

Oct 10, 2015

1 PARCv1 Instruction Cache

1.a Categorizing Cache Misses

Addr	Instruction	Iteration 1	Iteration 2
loop:			
0x108	addiu r1, r1, -1	compulsory	
0x10c	addiu r2, r2, -1		
0x110	j foo	compulsory	conflict
...			
foo:			
0x218	addiu r6, r6, 1	compulsory	conflict
0x21c	bne r1, r0, loop		

Figure 1: Cache Miss Type

1.b Average Memory Access Latency

Looking at iteration 2, we can see that there are 2 misses out of the 5 instructions. Therefore the miss rate for 64 iterations of the loop is 0.4. The average memory access latency is: $AMAL = (\text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Penalty})$ $AMAL = 1 + (0.4 \times 5)$ $AMAL = 3$ cycles The AMAL is dominated by conflict misses, as shown by the miss chart above. Compulsory misses only occur on the first iteration of the loop.

1.c Set-Associativity

The cache performance will increase significantly, because there will no longer be conflict misses during the loop. With this new cache microarchitecture, only compulsory misses will be left.

2 Page-Based Memory Translation

2.a Two-Level Page Tables

The 16-bit virtual address is used as the following:

Bits 14-15	Bits 12-13	Bits 0-11
XX	XX	XXXXXXXXXXXX
Virtual Page Number		Page Offset
L1 Index	L2 Index	Page Offset

Figure 2: Virtual Address Usage

Page Tables:

Paddr of PTE	Page-Table Entry	
	Valid	Paddr
0xffffc	1	0xfffe0
0xffff8		
0xffff4		
0xffff0	1	0xfffb0
0xfffec	1	0x05000
0xfffe8	1	0x07000
0xfffe4		
0xfffe0		
0xfffdc		
0xfffd8		
0xfffd4		
0xfffd0		
0xfffcc		
0xfffc8		
0xfffc4		
0xfffc0		
0xfffb0	1	0x01000
0xfffb8	1	0x04000
0xfffb4	1	0x00000
0xfffb0		

Figure 3: Contents of Physical Memory with Page Tables

2.b Translation-Lookaside Buffer

Transaction Address	VPN	Page Offset	m/h	Total Num Mem Accesses	Way 0		Way 1	
					VPN	PPN	VPN	PPN
0xeff4	0xe	0xff4	m	3	-	-	-	-
0x2ff0	0x2	0xff0	m	3	0xe	0x07		
0xeff8	0xe	0xff8	h	1			0x2	0x04
0x2ff4	0x2	0xff4	h	1				
0xeffc	0xe	0xffc	h	1				
0x2ff8	0x2	0xff8	h	1				
0xf000	0xf	0x000	m	3				
0x2ffc	0x2	0xffc	h	1	0xf	0x05		
0xf004	0xf	0x004	h	1				
0x3000	0x3	0x000	m	3				
0xf008	0xf	0x008	h	1			0x3	0x01
0x3004	0x3	0x004	h	1			0x2	0x04
0xf00c	0xf	0x00c	h	1				
0x3008	0x3	0x008	h	1				

Figure 4: TLB Contents Over Time

3 Impact of Cache Access Time and Replacement Policy

3.a Miss Rate Analysis

Transaction												
Address	tag	idx	m/h	L0	L1	L2	L3	L4	L5	L6	L7	
0x024	0x0	0x2	m	-	-	-	-	-	-	-	-	-
0x030	0x0	0x3	m			0x0						
0x07c	0x0	0x7	m				0x0					
0x070	0x0	0x7	h								0x0	
0x100	0x2	0x0	m									
0x110	0x2	0x1	m	0x2								
0x204	0x4	0x0	m		0x2							
0x214	0x4	0x1	m	0x4								
0x308	0x6	0x0	m		0x4							
0x110	0x2	0x1	m	0x6								
0x114	0x2	0x1	h		0x2							
0x118	0x2	0x1	h									
0x11c	0x2	0x1	h									
0x410	0x8	0x1	m									
0x110	0x2	0x1	m		0x8							
0x510	0xa	0x1	m		0x2							
0x110	0x2	0x1	m		0xa							
0x610	0xc	0x1	m		0x2							
0x110	0x2	0x1	m		0xc							
0x710	0xe	0x1	m		0x2							
Number of Misses = 16												
Miss Rate = 0.8												

Figure 5: Direct-Mapped Cache Contents Over Time

Transaction				Set 0		Set 1		Set 2		Set 3	
Address	tag	idx	m/h	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1
0x024	0x0	0x2	m	-	-	-	-	-	-	-	-
0x030	0x0	0x3	m					0x0			
0x07c	0x1	0x3	m							0x0	
0x070	0x1	0x3	h								0x1
0x100	0x4	0x0	m								
0x110	0x4	0x1	m	0x4							
0x204	0x8	0x0	m			0x4					
0x214	0x8	0x1	m		0x8						
0x308	0xc	0x0	m				0x8				
0x110	0x4	0x1	h	0xc							
0x114	0x4	0x1	h								
0x118	0x4	0x1	h								
0x11c	0x4	0x1	h								
0x410	0x10	0x1	m								
0x110	0x4	0x1	h				0x10				
0x510	0x14	0x1	m								
0x110	0x4	0x1	h				0x14				
0x610	0x18	0x1	m								
0x110	0x4	0x1	h				0x18				
0x710	0x1c	0x1	m								
Number of Misses = 12											
Miss Rate = 0.6											

Figure 6: Two-Way Set-Associative Cache Contents Over Time with LRU Replacement

Transaction				Set 0		Set 1		Set 2		Set 3	
Address	tag	idx	m/h	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1
0x024	0x0	0x2	m	-	-	-	-	-	-	-	-
0x030	0x0	0x3	m					0x0			
0x07c	0x1	0x3	m							0x0	
0x070	0x1	0x3	h								0x1
0x100	0x4	0x0	m								
0x110	0x4	0x1	m	0x4							
0x204	0x8	0x0	m			0x4					
0x214	0x8	0x1	m		0x8						
0x308	0xc	0x0	m				0x8				
0x110	0x4	0x1	h	0xc							
0x114	0x4	0x1	h								
0x118	0x4	0x1	h								
0x11c	0x4	0x1	h								
0x410	0x10	0x1	m								
0x110	0x4	0x1	h				0x10				
0x510	0x14	0x1	m								
0x110	0x4	0x1	m			0x14					
0x610	0x18	0x1	m								
0x110	0x4	0x1	m				0x18				
0x710	0x1c	0x1	m								
Number of Misses = 14											
Miss Rate = 0.7											

Figure 7: Two-Way Set-Associative Cache Contents Over Time with FIFO Replacement

3.b Sequential Tag Check then Memory Access

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
tag_decoder	$3 + 2 \times 2$	7
tag_mem	$10 + \lceil (4+27)/16 \rceil$	12
tag_cmp	$3 + 2\lceil \log_2(26) \rceil$	13
tag_and	$2 - 1$	1
data_decoder	$3 + 2 \times 3$	9
data_mem	$10 + \lceil (8+128)/16 \rceil$	19
rdata_mux	$3\lceil \log_2(4) \rceil + \lceil 32/8 \rceil$	10
rdata_reg_M1	1	1
Total		73
addr_reg_M0	1	1
tag_decoder	$3 + 2 \times 2$	7
tag_mem	$10 + \lceil (4+27)/16 \rceil$	12
tag_cmp	$3 + 2\lceil \log_2(26) \rceil$	13
tag_and	$2 - 1$	1
data_decoder	$3 + 2 \times 3$	9
data_mem	$10 + \lceil (8+128)/16 \rceil$	19
Total		62

Figure 8: Critical Path and Cycle Time for 2-Way Set-Associative Cache with Serialized Tag Check before Data Access

The reason that the 2-way set-associative microarchitecture is slower than the direct-mapped microarchitecture is the need for the tag check result to go through the data_decoder. It happens that the data_decoder's delay is relatively significant (9τ). This connection is needed so that the data can be outputted from the correct way.

3.c Parallel Read Hit Path

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
addr_mux	$3\lceil \log_2(2) \rceil + \lceil 5/8 \rceil$	4
data_decoder	$3 + 2 \times 3$	9
data_mem	$10 + \lceil (8+128)/16 \rceil$	19
rdata_mux	$3\lceil \log_2(4) \rceil + \lceil 32/8 \rceil$	10
rdata_reg_M1	1	1
Total		44

Figure 9: Critical Path and Cycle Time for Direct Mapped Cache with Parallel Read Hit

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
addr_mux	$3\lceil \log_2(2) \rceil + \lceil 5/8 \rceil$	4
data_decoder	$3 + 2 \times 2$	7
data_mem	$10 + \lceil (8+128)/16 \rceil$	19
rdata_mux	$3\lceil \log_2(4) \rceil + \lceil 32/8 \rceil$	10
way_mux	$3\lceil \log_2(2) \rceil + \lceil 32/8 \rceil$	7
rdata_reg_M1	1	1
Total		49

Figure 10: Critical Path and Cycle Time for 2-Way Set-Associative Cache with Parallel Read Hit

The reason that the 2-way set-associative microarchitecture is slower than the direct-mapped microarchitecture is the way_mux, which is needed to output the data from the correct way. This mux has a delay of 7τ , which is relatively significant.

3.d Pipelined Write Hit Path

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
tag_decoder	$3 + 2 \times 3$	9
tag_mem	$10 + \lceil (8+26)/16 \rceil$	13
tag_cmp	$3 + 2\lceil \log_2(25) \rceil$	13
tag_and	$2 - 1$	1
wen_and	$2 - 1$	1
wen_reg_M1	1	1
Total		39

Figure 11: Critical Path and Cycle Time for Direct Mapped Cache with Pipelined Write Hit

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
tag_decoder	$3 + 2 \times 2$	7
tag_mem	$10 + \lceil (4+27)/16 \rceil$	12
tag_cmp	$3 + 2\lceil \log_2(25) \rceil$	13
tag_and	$2 - 1$	1
wen_and	$2 - 1$	1
wen_reg_M1	1	1
Total		36

Figure 12: Critical Path and Cycle Time for 2-Way Set-Associative Cache with Pipelined Write Hit