

ECE 4750 PSET 2

Tim Yao (ty252)

Oct 10, 2015

1 PARCv1 Instruction Cache

1.a Categorizing Cache Misses

Addr	Instruction	Iteration 1	Iteration 2
loop:			
0x108	addiu r1, r1, -1	compulsory	
0x10c	addiu r2, r2, 1		
0x110	j foo	compulsory	conflict
...			
foo:			
0x218	addiu r6, r6, 1	compulsory	conflict
0x21c	bne r1, r0, loop		

Figure 1: Cache Miss Type

1.b Average Memory Access Latency

Looking at iteration 2, we can see that there are 2 misses out of the 5 instructions. Therefore the miss rate for 64 iterations of the loop is 0.4.

The average memory access latency is:

$$\text{AMAL} = (\text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Penalty})$$

$$\text{AMAL} = 1 + (0.4 \times 5)$$

$$\text{AMAL} = 3 \text{ cycles}$$

The AMAL is dominated by conflict misses, as shown by the miss chart above. Compulsory misses only occur on the first iteration of the loop.

1.c Set-Associativity

The cache performance will increase significantly, because there will no longer be conflict misses during the loop. With this new cache microarchitecture, only compulsory misses will be left.

2 Page-Based Memory Translation

2.a Two-Level Page Tables

The 16-bit virtual address is used as the following:

Bits 14-15	Bits 12-13	Bits 0-11
XX	XX	XXXXXXXXXXXXXX
Virtual Page Number		Page Offset
L1 Index	L2 Index	Page Offset

Figure 2: Virtual Address Usage

Page Tables:

Paddr of PTE	Page-Table Entry	
	Valid	Paddr
0xffffc	1	0xfffe0
0xffff8		
0xffff4		
0xffff0	1	0xfffb0
0xfffec	1	0x05000
0xfffe8	1	0x07000
0xfffe4		
0xfffe0		
0xfffdc		
0xfffd8		
0xfffd4		
0xfffd0		
0xfffcc		
0xfffc8		
0xfffc4		
0xfffc0		
0xfffb0	1	0x01000
0xfffb8	1	0x04000
0xfffb4	1	0x00000
0xfffb0		

Figure 3: Contents of Physical Memory with Page Tables

2.b Translation-Lookaside Buffer

Transaction Address	VPN	Page Offset	m/h	Total Num Mem Accesses	TLB Way 0		TLB Way 1	
					VPN	PPN	VPN	PPN
0xeff4	0xe	0xff4	m	3	-	-	-	-
0x2ff0	0x2	0xff0	m	3	0xe	0x07	-	-
0xeff8	0xe	0xff8	h	1	0xe	0x07	0x2	0x04
0x2ff4	0x2	0xff4	h	1	0xe	0x07	0x2	0x04
0xeffc	0xe	0xffc	h	1	0xe	0x07	0x2	0x04
0x2ff8	0x2	0xff8	h	1	0xe	0x07	0x2	0x04
0xf000	0xf	0x000	m	3	0xe	0x07	0x2	0x04
0x2ffc	0x2	0xffc	h	1	0xf	0x05	0x2	0x04
0xf004	0xf	0x004	h	1	0xf	0x05	0x2	0x04
0x3000	0x3	0x000	m	3	0xf	0x05	0x2	0x04
0xf008	0xf	0x008	h	1	0xf	0x05	0x3	0x01
0x3004	0x3	0x004	h	1	0xf	0x05	0x2	0x04
0xf00c	0xf	0x00c	h	1	0xf	0x05	0x2	0x04
0x3008	0x3	0x008	h	1	0xf	0x05	0x2	0x04

Figure 4: TLB Contents Over Time

3 Impact of Cache Access Time and Replacement Policy

3.a Miss Rate Analysis

Transaction				L0	L1	L2	L3	L4	L5	L6	L7
Address	tag	idx	m/h								
0x024	0x0	0x2	m	-	-	-	-	-	-	-	-
0x030	0x0	0x3	m			0x0					
0x07c	0x0	0x7	m				0x0				
0x070	0x0	0x7	h								0x0
0x100	0x2	0x0	m								
0x110	0x2	0x1	m	0x2							
0x204	0x4	0x0	m		0x2						
0x214	0x4	0x1	m	0x4							
0x308	0x6	0x0	m		0x4						
0x110	0x2	0x1	m	0x6							
0x114	0x2	0x1	h		0x2						
0x118	0x2	0x1	h								
0x11c	0x2	0x1	h								
0x410	0x8	0x1	m								
0x110	0x2	0x1	m		0x8						
0x510	0xa	0x1	m		0x2						
0x110	0x2	0x1	m		0xa						
0x610	0xc	0x1	m		0x2						
0x110	0x2	0x1	m		0xc						
0x710	0xe	0x1	m		0x2						
Number of Misses = 16											
Miss Rate = 0.8											

Figure 5: Direct-Mapped Cache Contents Over Time

Transaction				Set 0		Set 1		Set 2		Set 3	
Address	tag	idx	m/h	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1
0x024	0x0	0x2	m	-	-	-	-	-	-	-	-
0x030	0x0	0x3	m					0x0			
0x07c	0x1	0x3	m							0x0	
0x070	0x1	0x3	h								0x1
0x100	0x4	0x0	m								
0x110	0x4	0x1	m	0x4							
0x204	0x8	0x0	m			0x4					
0x214	0x8	0x1	m		0x8						
0x308	0xc	0x0	m				0x8				
0x110	0x4	0x1	h	0xc							
0x114	0x4	0x1	h								
0x118	0x4	0x1	h								
0x11c	0x4	0x1	h								
0x410	0x10	0x1	m								
0x110	0x4	0x1	h				0x10				
0x510	0x14	0x1	m								
0x110	0x4	0x1	h				0x14				
0x610	0x18	0x1	m				0x18				
0x110	0x4	0x1	h								
0x710	0x1c	0x1	m								
Number of Misses = 12											
Miss Rate = 0.6											

Figure 6: Two-Way Set-Associative Cache Contents Over Time with LRU Replacement

Transaction				Set 0		Set 1		Set 2		Set 3	
Address	tag	idx	m/h	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1	Way 0	Way 1
0x024	0x0	0x2	m	-	-	-	-	-	-	-	-
0x030	0x0	0x3	m					0x0			
0x07c	0x1	0x3	m							0x0	
0x070	0x1	0x3	h								0x1
0x100	0x4	0x0	m								
0x110	0x4	0x1	m	0x4							
0x204	0x8	0x0	m			0x4					
0x214	0x8	0x1	m		0x8						
0x308	0xc	0x0	m				0x8				
0x110	0x4	0x1	h	0xc							
0x114	0x4	0x1	h			0x4					
0x118	0x4	0x1	h								
0x11c	0x4	0x1	h								
0x410	0x10	0x1	m								
0x110	0x4	0x1	h				0x10				
0x510	0x14	0x1	m								
0x110	0x4	0x1	m			0x14					
0x610	0x18	0x1	m								
0x110	0x4	0x1	m				0x18				
0x710	0x1c	0x1	m								
Number of Misses = 14											
Miss Rate = 0.7											

Figure 7: Two-Way Set-Associative Cache Contents Over Time with FIFO Replacement

3.b Sequential Tag Check then Memory Access

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
tag_decoder	$3+2\times 2$	7
tag_mem	$10 + \lfloor (4+27)/16 \rfloor$	12
tag_cmp	$3+2\lceil \log_2(26) \rceil$	13
tag_and	2-1	1
data_decoder	$3+2\times 3$	9
data_mem	$10 + \lfloor (8+128)/16 \rfloor$	19
rdata_mux	$3\lceil \log_2(4) \rceil + \lfloor 32/8 \rfloor$	10
rdata_reg_M1	1	1
Total		73
addr_reg_M0	1	1
tag_decoder	$3+2\times 2$	7
tag_mem	$10 + \lfloor (4+27)/16 \rfloor$	12
tag_cmp	$3+2\lceil \log_2(26) \rceil$	13
tag_and	2-1	1
data_decoder	$3+2\times 3$	9
data_mem	$10 + \lfloor (8+128)/16 \rfloor$	19
Total		62

Figure 8: Critical Path and Cycle Time for 2-Way Set-Associative Cache with Serialized Tag Check before Data Access

The reason that the 2-way set-associative microarchitecture is slower than the direct-mapped microarchitecture is the need for the tag check result to go through the data_decoder. It happens that the data_decoder's delay is relatively significant (9τ).

3.c Parallel Read Hit Path

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
addr_mux	$3\lceil \log_2(2) \rceil + \lfloor 5/8 \rfloor$	4
data_decoder	$3+2\times 3$	9
data_mem	$10 + \lfloor (8+128)/16 \rfloor$	19
rdata_mux	$3\lceil \log_2(4) \rceil + \lfloor 32/8 \rfloor$	10
rdata_reg_M1	1	1
Total		44

Figure 9: Critical Path and Cycle Time for Direct Mapped Cache with Parallel Read Hit

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
addr_mux	$3\lceil \log_2(2) \rceil + \lfloor 4/8 \rfloor$	4
data_decoder	$3+2\times 2$	7
data_mem	$10 + \lfloor (4+128)/16 \rfloor$	19
rdata_mux	$3\lceil \log_2(4) \rceil + \lfloor 32/8 \rfloor$	10
way_mux	$3\lceil \log_2(2) \rceil + \lfloor 32/8 \rfloor$	7
rdata_reg_M1	1	1
Total		49

Figure 10: Critical Path and Cycle Time for 2-Way Set-Associative Cache with Parallel Read Hit

INSERT SOME SHIT HERE

3.d Pipelined Write Hit Path

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
tag_decoder	$3+2\times 3$	9
tag_mem	$10 + \lceil (8+26)/16 \rceil$	13
tag_cmp	$3+2\lceil \log_2(25) \rceil$	13
tag_and	2-1	1
wen_and	2-1	1
wen_reg_M1	1	1
Total		39

Figure 11: Critical Path and Cycle Time for Direct Mapped Cache with Pipelined Write Hit

Component	Delay Equation	Delay(τ)
addr_reg_M0	1	1
tag_decoder	$3+2\times 2$	7
tag_mem	$10 + \lceil (4+27)/16 \rceil$	12
tag_cmp	$3+2\lceil \log_2(26) \rceil$	13
tag_and	2-1	1
wen_and	2-1	1
wen_reg_M1	1	1
Total		36

Figure 12: Critical Path and Cycle Time for 2-Way Set-Associative Cache with Pipelined Write Hit

INSERT SOME SHIT HERE

3.e Average Memory Access Latency

4 Two-Cycle Pipelined Integer ALU and Multiplier

4.a Part 4.A Control and Data Hazard Latencies

Jump Resolution Latency = 2

Branch Resolution Latency = 4

ALU-use Delay Latency = 2

Load-use Delay Latency = 3

4.b Resolving Data Hazards with Software Scheduling

```

1      bne r1, r0, done
2      lw  r5, 0(r2)
3      lw  r6, 0(r3)
4      addiu r8, r4, 4
5      nop
6      addu r7, r5, r6
7      nop
8      sw  r7, 0(r8)
9      ...
10     done: addiu r10, r9, 1

```

By swapping the addu and addiu instructions, we can avoid 1 nop for the dependence between the addu and the lw instructions. Regardless of how the addiu and addu instructions are scheduled, there will always be a nop right before the sw instruction due to the dependence.

Dynamic Transaction			Cycle												
			1	2	3	4	5	6	7	8	9	10	11	12	13
1	bne	r1, r0, done	F	D	X0	X1	M	W							
2	lw	r5, 0(r2)		F	D	X0	X1	M	W						
3	lw	r6, 0(r3)			F	D	X0	X1	M	W					
4	addiu	r8, r4, 4				F	D	X0	X1	M	W				
5	nop						F	D	X0	X1	M	W			
6	addu	r7, r5, r6						F	D	X0	X1	M	W		
7	nop								F	D	X0	X1	M	W	
8	sw	r7, 0(r8)								F	D	X0	X1	M	W

Figure 13: Software Scheduling Pipeline

Instruction 6 (addu)'s D stage depends on instruction 2 (lw)'s M stage and instruction 3 (lw)'s M stage. Instruction 7 (sw)'s D stage depends on instruction 4 (addiu)'s X1 stage and instruction 6 (addu)'s X1 stage.

4.c Resolving Data Hazards with Stalling

Dynamic Transaction			Cycle													
			1	2	3	4	5	6	7	8	9	10	11	12	13	14
1	bne	r1, r0, done	F	D	X0	X1	M	W								
2	lw	r5, 0(r2)		F	D	X0	X1	M	W							
3	lw	r6, 0(r3)			F	D	X0	X1	M	W						
4	addu	r7, r5, r6				F	D	D	D	X0	X1	M	W			
5	addiu	r8, r4, 4					F	F	F	D	X0	X1	M	W		
6	sw	r7, 0(r8)								F	F	D	X0	X1	M	W

Figure 14: Hardware Stalling Pipeline

Instruction 6 (addu)'s D stage depends on instruction 2 (lw)'s M stage and instruction 3 (lw)'s M stage.
 Instruction 7 (sw)'s D stage depends on instruction 4 (addiu)'s X1 stage and instruction 6 (addu)'s X1 stage.

4.d New Stall Signal

```

1 ostall_load_use_X_rs_D
2   =   val_D && rs_en_D && (val_X0 || val_X1) && (rf_wen_X0 || rf_wen_X1)
3     && ((inst_rs_D == rf_waddr_X0) || (inst_rs_D == rf_waddr_X1))
4     && ((rf_waddr_X0 != 0) || (rf_waddr_X1 != 0))
5     && ((op_X0 == LW) || (op_X1 == LW))
6 ostall_load_use_X_rt_D
7   =   val_D && rt_en_D && (val_X0 || val_X1) && (rf_wen_X0 || rf_wen_X1)
8     && ((inst_rt_D == rf_waddr_X0) || (inst_rt_D == rf_waddr_X1))
9     && ((rf_waddr_X0 != 0) || (rf_waddr_X1 != 0))
10    && ((op_X0 == LW) || (op_X1 == LW))
11 ostall_alu_use_X_rs_D
12   =   val_D && rs_en_D && val_X0 && rf_wen_X0
13     && (inst_rs_D == rf_waddr_X0)
14     && (rf_waddr_X0 != 0)
15 ostall_alu_use_X_rt_D
16   =   val_D && rt_en_D && val_X0 && rf_wen_X0
17     && (inst_rt_D == rf_waddr_X0)
18     && (rf_waddr_X0 != 0)
19 stall_D
20   =   val_D && !squash_D
21     && ((ostall_load_use_X_rs_D || ostall_load_use_X_rt_D) ||
22         (ostall_alu_use_X_rs_D || ostall_alu_use_X_rt_D))

```

4.e Resolving Control Hazards with Scheduling and Speculation

Dynamic Transaction	Cycle									
	1	2	3	4	5	6	7	8	9	10
1 bne r1, r0, done	F	D	X0	X1	M	W				
2 lw r5, 0(r2)		F	D	X0	-	-	-			
3 lw r6, 0(r3)			F	D	-	-	-	-		
4 addu r7, r5, r6				F	-	-	-	-	-	
5 addiu r10, r9, 1					F	D	X0	X1	M	W

Figure 15: Squashing Pipeline

Dynamic Transaction	Cycle									
	1	2	3	4	5	6	7	8	9	10
1 bne r1, r0, done	F	D	X0	X1	M	W				
2 lw r5, 0(r2)		F	D	X0	X1	M	W			
3 lw r6, 0(r3)			F	D	-	-	-	-		
4 addu r7, r5, r6				F	-	-	-	-	-	
5 addiu r10, r9, 1					F	D	X0	X1	M	W

Figure 16: Branch Delay Slot Pipeline